


Compositionality Entails Sequentializability

Pranav Garg, P. Madhusudan

University of Illinois at Urbana-Champaign

What is Sequentializability ?



- Sequentialization requires that the size of sequential program is of the same order as the size of concurrent program.
 - The constant in $O(N)$ and $O(L)$ must be independent of #threads.
- For non-recursive concurrent programs, global simulation is always possible to obtain a sequential program.
 - Leads to a state-space explosion (L^n where n is #threads)
 - Is not a sequentialization.
- Recursive programs are even harder.

Why are sequentializations appealing ?

Practically

- Allows the use of analysis tools developed for sequential programs to be used directly for concurrent programs.
 - Deductive verification
 - Symbolic model checking
 - Predicate abstraction
 - Symbolic test case generation

Theoretically

- Intriguing
 - When are concurrent programs sequentializable?

Earlier Results

- Analysis of concurrent programs under a bounded number of context-rounds (no thread creation).
 - Lal and Reps [CAV 2008].
 - Used for bounded model checking of concurrent programs.
 - STORM [Lahiri, Qadeer and Rakamaric] and Poirot from Microsoft.
 - LaTorre, Madhusudan and Parlato [CAV 2009]
 - Shown to be more efficient for explicit model checking with state caching.
 - LaTorre, Madhusudan and Parlato [Unpublished]
 - Parameterized programs with unbounded number of threads but a bounded number of rounds.
- Analysis of concurrent programs with dynamic thread creation under a delay bound.
 - Emmi, Qadeer and Rakamaric [POPL 2011]

Compositionality entails sequentializability

Compositional semantics of a concurrent program with respect to a set of auxiliary variables can be sequentialized.

- It generalizes the prior sequentializations known for the under-approximate context-bounded analysis.
 - Bounding the number of context switches makes them amenable to compositional reasoning.
- Compositional semantics: Over-approximation
 - Can be used to “prove” concurrent programs correct.

Overview

- Jones style rely-guarantee proofs
- Compositional semantics for concurrent program
- Main theorem, intuition behind the sequentialization.
- Experimental results.
- Conclusion.

Rely-Guarantee Proofs

Hoare style method for proving concurrent program.

$$P \models (pre, post, rely, guar)$$

- $pre, post$ are unary predicates defining subsets of states.
- $rely, guar$ are binary relations defining transformations to the shared state.

Parallel compositional rule:

$$\frac{\begin{array}{ll} guar_1 \Rightarrow rely_2, & guar_2 \Rightarrow rely_1, \\ P \models (pre, post_1, rely_1, guar_1), & Q \models (pre, post_2, rely_2, guar_2) \\ rely \Rightarrow rely_1, rely \Rightarrow rely_2, & (guar_1 \vee guar_2) \Rightarrow guar \end{array}}{P || Q \models (pre, post_1 \wedge post_2, rely, guar)}$$

Auxiliary variables

- Parallel compositional rule in itself is not complete.

- Example:

```
pre:  x = 0

      P1      ||      P2
atomic {      atomic {
  x := x + 1;  x := x + 1;
}             }

post:  x = 2
```

- Impossible to come up with consistent *rely-guar* conditions which are strong enough to prove the *post*.

Auxiliary variables

pre: $x = 0 \wedge pc_1 = 0 \wedge pc_2 = 0$

P_1 $\text{atomic } \{$ $x := x + 1;$ <u>$pc_1 := 1;$</u> $\}$		P_2 $\text{atomic } \{$ $x := x + 1;$ <u>$pc_2 := 1;$</u> $\}$
<i>post:</i> $x = 2$		

Auxiliary variables

- Since auxiliary variables are not read, the semantics of the concurrent program remains unchanged.
- Auxiliary variables are the local variables which are required to be exposed to the environment to prove the concurrent program compositionally.

Compositional Semantics wrt. Auxiliary variables

$$P = P_1 \parallel P_2$$

- L_1, L_2 : the set of local variables of the individual threads.
- S : the set of shared variables.
- $A \subseteq L_1 \cup L_2$: the set of auxiliary variables.
- δ_1, δ_2 : local (binary) transition relations of P_1 and P_2

Compositional Semantics is defined by four sets:

$$\begin{array}{ll} R_1 \subseteq (Val_{L_1} \times Val_S \times Val_{A \cap L_2}), & R_i : \text{Reachable state in } P_i \\ R_2 \subseteq (Val_{L_2} \times Val_S \times Val_{A \cap L_1}), & \\ Guar_1, Guar_2 \subseteq (Val_S \times Val_A \times Val_S \times Val_A), & Guar_i : \text{Guarantee that } P_i \text{ promises.} \end{array}$$

Compositional Semantics wrt. Auxiliary variables

a) Initialization:

- R_i contains the set $\{(l_i, s, t) \mid l_i \cup s \cup t \in \text{Init} \downarrow (L_i \cup A \cup S)\}$.

b) Transitions of P_1 : If $(l_1, s, t) \in R_1$ and $\delta_1(l_1, s, l'_1, s')$ holds, then

- **Local update:** $(l'_1, s', t) \in R_1$.
- **Update to guarantee:** $(s, l_1 \downarrow A \cup t, s', l'_1 \downarrow A \cup t) \in \text{Guar}_1$.

c) Transitions of P_2 : Similar to (b)

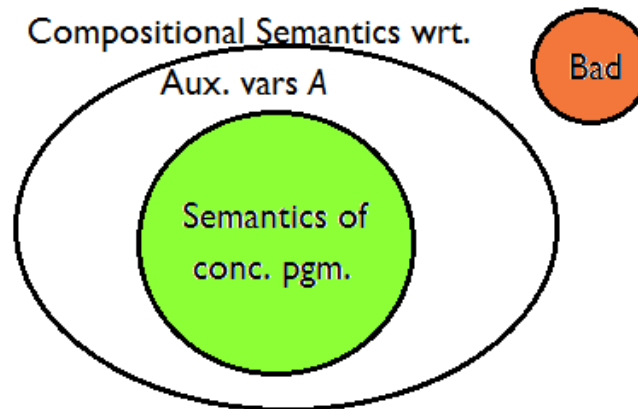
d) Interference:

- If $(l_1, s, t) \in R_1$ and $(s, l_1 \downarrow A \cup t, s', t') \in \text{Guar}_2$, then $(l_1, s', t' \downarrow L_2) \in R_1$.
- If $(l_2, s, t) \in R_2$ and $(s, l_2 \downarrow A \cup t, s', t') \in \text{Guar}_1$, then $(l_2, s', t' \downarrow L_1) \in R_2$.

$$\text{Reach}_A = \{(l_1, s, l_2) \mid (l_1, s, l_2 \downarrow A) \in R_1 \text{ and } (l_2, s, l_1 \downarrow A) \in R_2\}$$

Compositional Semantics wrt. Auxiliary variables

- Existence of a Jones style rely-guarantee proof →
Compositional semantics of the concurrent program with respect to the corresponding auxiliary variables is correct.
- Note:
 - Compositional semantics is an over-approximation of the standard operational semantics of the concurrent programs.



The main result

- Given a concurrent program C with auxiliary variables A , we can build a sequential program $S_{C,A}$ such that
 - The compositional semantics of the concurrent program with respect to the auxiliary variables A is correct iff $S_{C,A}$ is correct.
 - At any point, the scope of $S_{C,A}$ keeps a constant number of copies of the variables of C .

Intuition behind the sequentialization

- Let there be methods G_1 and G_2 which semantically capture the guarantees $Guar_1$ and $Guar_2$ respectively.

$$G_i(\hat{s}, \hat{a}) = \{(s, a) \mid (\hat{s}, \hat{a}, s, a) \in Guar_i\}$$

Once we have G_2 , we can compute the reachable states R_1 according to:


```
while(*) {  
  if (*) then  
    <<simulate a transition of P1>>  
  else  
    (s,a) := G2(s,a);  
  fi  
}
```

- Track the local state L_1 , shared state and the auxiliary state of the second thread
- On two successive calls to G_2 , there is no preservation of the local state of P_2 .

Intuition behind the sequentialization

The method G_2 (and similarly G_1) can be implemented as:

```
G2(s#, a#) {  
  <<initialize variables of P2>>  
  while(*) {  
    if (*) then  
      <<simulate a transition of P2>>  
    else  
      (s,a) := G1(s,a);  
    fi  
  }  
  assume (local and shared state is consistent with s#, a#);  
  while(*) {  
    <<simulate a transition of P2>>  
  }  
  return (s,a);  
}
```



Code for R_2

An example sequentialization

<pre> decl int x, pc1, pc2; decl int x*, pc1*, pc2*; decl bool z, term; main begin G_1(); return end I_i: if(term = true) then return fi if(!z & *) then x', pc1', pc2' := x*, pc1*, pc2*; G_{3-i}(); z, term:=false, false; x*, pc1*, pc2* := x', pc1', pc2' fi if(x = x* & pc1 = pc1* & pc2 = pc2* & *) then z := true fi if(z & *) then term := true; return fi </pre>	<pre> G_1() begin z, term:=false, false; x*, pc1*, pc2* := x, pc1, pc2 x, pc1, pc2 := 0, 0, 0; main_1(); assume(term = true); return end main_1() begin decl int x', pc1', pc2'; I_1 x := x + 1; pc1 := 1; I_1 assert(spec_1); I_1 return end </pre>	<pre> G_2() begin z, term:=false, false; x*, pc1*, pc2*:= x, pc1, pc2 x, pc1, pc2:=0, 0, 0; main_2(); assume(term = true); return end main_2() begin decl int x', pc1', pc2'; I_2 x := x + 2; pc2 := 1; I_2 assert(spec_2); I_1 return end </pre>
--	---	--

Applications of the theorem

- Deductive verification
- Predicate abstraction

Deductive Verification

Concurrent program C with
Auxiliary variables A

➔ Sequentialization $S_{C,A}$

Jones style
Rely/guarantee annotation

➔ Hoare-style pre-post
conditions & assertions

- *rely/guar*

➔ summaries of G_1 and G_2

- *pre/post*

➔ pre/post conditions

- assertions

➔ assertions

- loop invariants

➔ loop invariants

- Use deductive verification tools (like Boogie/Z3) to verify the correctness of $S_{C,A}$ and hence the correctness of the rely/guarantee proof of C .

Deductive Verification

Experimental results

Programs	Concurrent pgm			Sequential pgm		Time
	#Threads	#LOC	#Lines of annotations	#LOC	#Lines of annotations	
X++	2	38	5	113	6	8s
Lock	2	50	9	184	10	122s
Peterson	2	52	35	232	36	145s
Bakery	2	55	8	147	13	18s
ArrayIndexSearch	2	74	17	222	21	126s
GCD	2	78	23	279	29	869s
Bluetooth	unbdd	69	20	276	55	107s

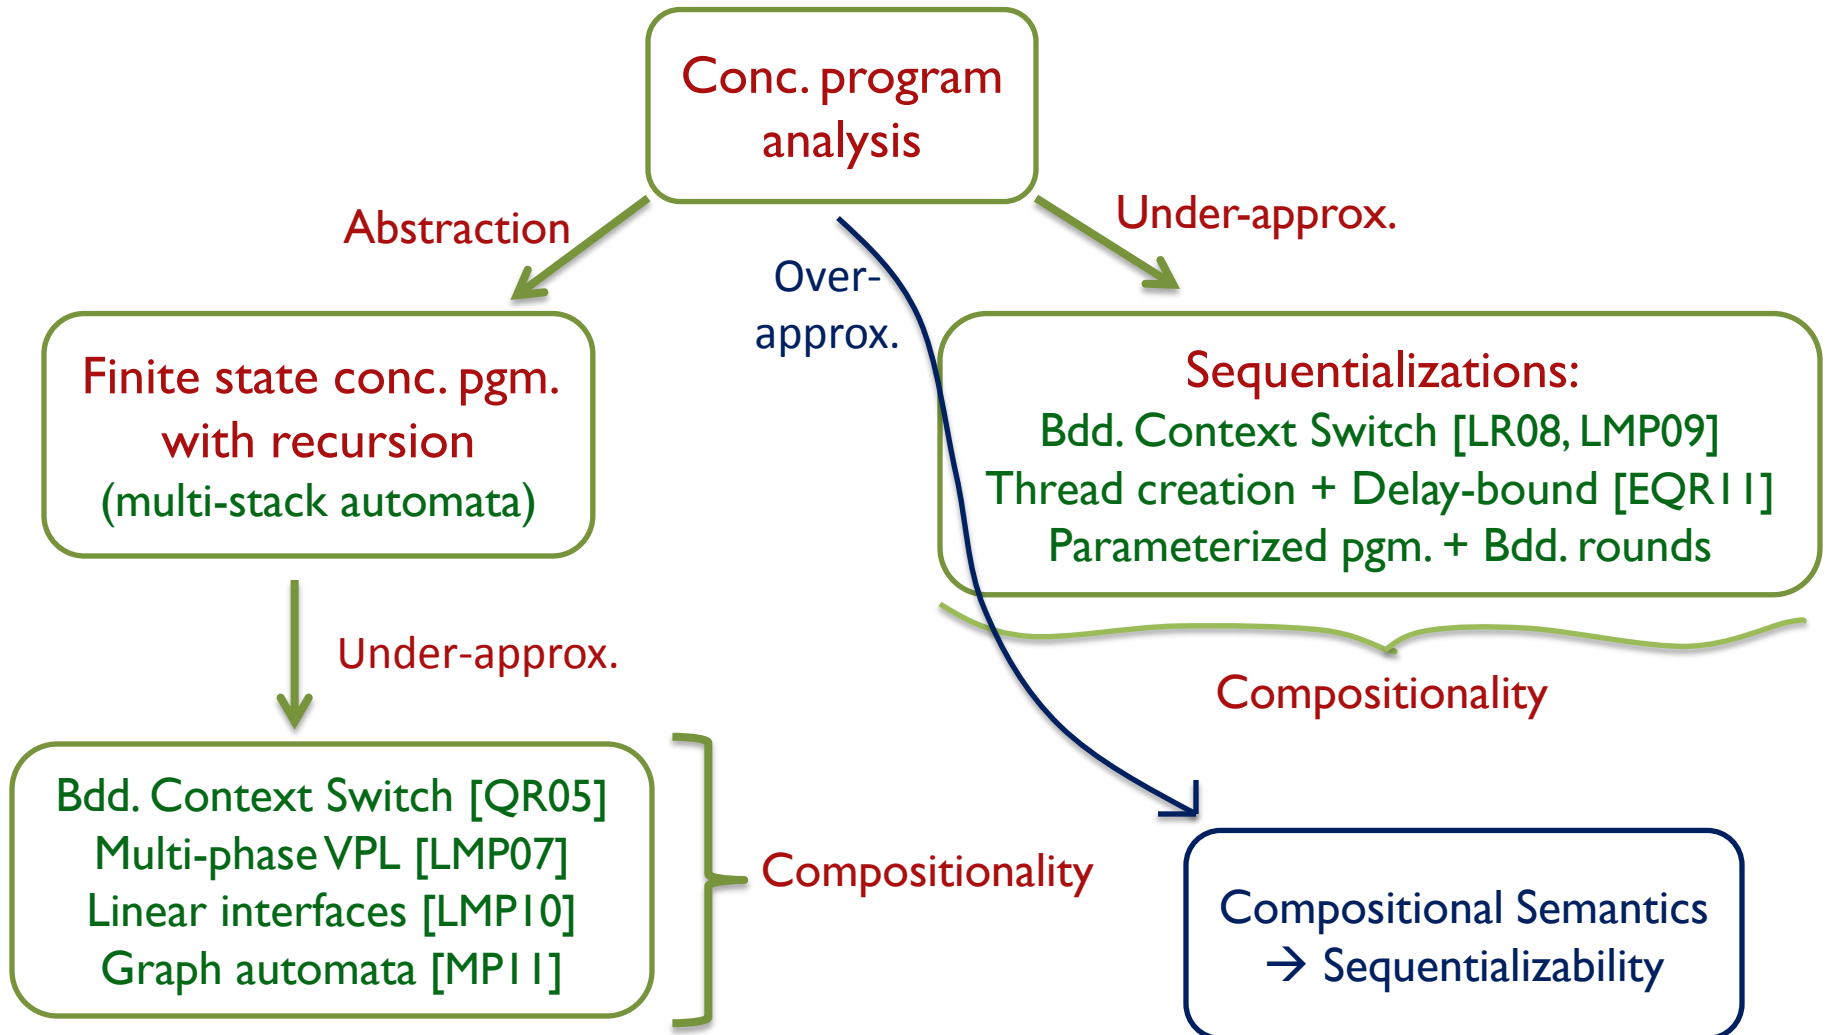
- Manually wrote Jones-style rely/guarantee annotations
- Boogie for Hoare-style verification of the corresponding sequential program.
- Experiments run on Intel dual-core with 1.6 GHz and 1 Gb RAM.

Predicate Abstraction

Concurrent program C + auxiliary variables A determined manually/heuristically \rightarrow Sequential program $S_{C,A}$

- Use a predicate abstraction tool (like SLAM) to prove $S_{C,A}$ and hence C correct.
 - The procedure is sound but incomplete.
 - If $S_{C,A}$ cannot be proved correct, cannot deduce anything (more auxiliary variables may be needed)
- We get a semi-automatic predicate abstraction tool for concurrent programs.
 - Determining auxiliary variables can be automated
[Cohen-Namjoshi, Gupta-Popeea-Rybalchenko]
- Used for proving programs: X++, Lock, Bakery and Peterson.

Broader context of the result



Conclusion

- Compositionality entails sequentializability.
- Can be used to prove concurrent programs correct.
- Our result could potentially have a number of applications:
 - Deductive verification
 - Predicate abstraction
 - Symbolic testing
 - Static analysis